
ECHOLIFE

Privacy-First Lifelong Digital Voice Identity
for People Facing Speech Loss

Complete Project Proposal

Team: **Gublet Gang**

Members: Srinjoy Pramanick · Soudip Biswas · Anwesha Das

College: JIS College of Engineering

Category: Software | Event: IGNIS 2k26 (IEM × UEM × Microsoft Student Society)

One-line pitch:

EchoLife is the world's first fully offline, privacy-first lifelong Digital Voice Identity system that lets people with progressive speech loss continue speaking in their own authentic voice — forever, and only on their own device.

This document presents the complete idea, deep technical approach, feasibility analysis, and impact assessment with special focus on the unique Encrypted On-Device Lifelong Digital Voice Identity feature.

1. Idea Title & Solution Overview

Idea Title

EchoLife – Privacy-First Lifelong Digital Voice Identity for Speech-Loss Patients

A. Proposed Solution

EchoLife is a privacy-first, fully on-device AI platform that captures a short voice sample from a person *before* speech loss occurs (ALS, stroke, Parkinson's, throat cancer, etc.), builds a permanent **Digital Voice Identity**, encrypts and stores it locally, and then allows the user to type (or speak residual speech) and instantly hear natural speech in *their own original voice* — completely offline, with zero cloud dependency.

It is not a generic text-to-speech app. It is a lifelong personal voice archive + real-time synthesizer that restores identity, not just communication.

B. Detailed Explanation of the Proposed Solution

- **Capture phase** (one-time, while the person can still speak): User records 10–30 seconds of natural speech via the mobile/web app.
- **Identity creation**: On-device models (Whisper for transcription + VALL-E-X-style zero-shot voice encoder) extract speaker embeddings, prosody, timbre, and speaking-style features to create a compact Digital Voice Identity profile.
- **Secure local storage**: The profile is encrypted (AES-256 + device-bound key) and stored in an on-device encrypted SQLite / secure enclave. No voice data ever leaves the device.
- **Daily use**: User types a message (or uses residual speech → Whisper on-device). The Edge AI inference engine (optimized via Qualcomm AI Hub / Snapdragon NPU) synthesizes speech in the preserved voice and plays it instantly.
- **Lifelong**: The same identity can be used for years, even after complete loss of speech. Optional future fine-tuning from residual vocalizations keeps the voice natural as the user's condition changes.

C. How it Addresses the Problem

Current solutions (cloud TTS, generic robotic voices, hospital voice banking that still relies on servers) destroy privacy and produce voices that do not sound like the person. EchoLife solves both problems simultaneously: **identity preservation + privacy + offline reliability**. Patients regain the ability to speak to family, doctors, and caregivers in their own voice, preserving dignity and emotional connection.

D. Innovation and Uniqueness

Core differentiators:

- Lifelong Digital Voice Identity (not one-shot cloning).
- 100% on-device Edge AI pipeline (no cloud at any stage after optional model download).
- Extremely short sample requirement + high naturalness.
- **Standout unique feature (deep-dived in Section 1.1 below): Encrypted On-Device Lifelong Digital Voice Identity with Continuous Edge Adaptation.**

1.1 Deep Dive – The Best Unique Feature

Encrypted On-Device Lifelong Digital Voice Identity with Continuous Edge Adaptation

This is the single feature that separates EchoLife from every existing voice-banking or TTS solution and the one on which judges will primarily evaluate depth of technical thinking.


Why existing solutions fall short

Most voice-cloning systems today are either cloud-only (ElevenLabs, PlayHT, Microsoft Azure, etc.) — creating a privacy nightmare for medical-grade voice data and becoming unusable without internet — or one-time offline cloning that produces a static model which quickly sounds robotic or drifts as the patient’s residual voice changes.

How EchoLife treats the voice as a living, encrypted digital organ

EchoLife builds a compact, encrypted Digital Voice Identity that lives entirely on the user’s device. The identity is device-bound: even if the phone is lost or stolen, the voice cannot be extracted without the user’s biometrics or device key. After the (optional) one-time model optimization via Qualcomm AI Hub, the entire pipeline — recognition, synthesis, and playback — runs offline on the Snapdragon NPU / equivalent edge accelerator. This makes the system usable in hospitals with no Wi-Fi, rural areas, or during network outages — critical real-world conditions for patients.

Architecture of the Unique Feature



Key technical strengths of this feature

- **Device-bound encryption:** Voice identity is protected by hardware-backed keys (Android Keystore / iOS Secure Enclave). Extraction is cryptographically hard without user presence.
- **Zero-shot + continuous adaptation path:** High-quality profile from only 10–30 seconds. Later residual speech or whispered attempts can gently fine-tune the embedding on-device (user-consented) so the voice does not “age out” or become mismatched.
- **True offline Edge AI:** After model quantization and NPU offload, no network is required. Works in ICUs, rural clinics, or during outages.
- **Tiny runtime footprint:** Quantized models + compact embedding make the solution practical on mid-range phones and future dedicated assistive hardware.
- **Identity, not just utility:** Restores the emotional and social dimension of hearing a loved one’s real voice — the core human need that generic TTS cannot meet.

This combination of lifelong identity + military-grade local privacy + true offline Edge AI does not exist in current commercial or research assistive tools. It is the feature that makes EchoLife fundamentally different.

2. Technical Approach

A. Technologies to be Used

Layer	Technologies
AI / ML	OpenAI Whisper (on-device ASR), VALL-E-X / equivalent zero-shot TTS encoder-decoder, custom voice-embedding models
Edge AI	Qualcomm AI Hub, Snapdragon X / Hexagon NPU optimization, ONNX Runtime / TensorRT Mobile
Development	Python, FastAPI (local backend), React.js / React Native (frontend)
Audio	PyAudio, Librosa, NumPy, torchaudio
Security & Storage	Encrypted SQLite, Android Keystore / iOS Secure Enclave, AES-256
Deployment	Fully offline mobile app + optional lightweight desktop companion

B. Methodology and Process of Implementation

Step-by-step working of the solution:

- **Record** – User records a short voice sample (10–30 s) through the app.
- **Preserve** – AI builds a Digital Voice Identity using Whisper + VALL-E-X-style encoder.
- **Store** – Voice identity is encrypted and stored securely on the device (Secure Enclave / Keystore).
- **Communicate** – User types a message (or speaks residual speech) whenever needed.
- **Generate** – Edge AI synthesizes natural speech in the user's preserved voice.
- **Output** – Natural speech is played instantly, even when completely offline.

Implementation Roadmap (Prototype Path)

- **Phase 1** – Desktop proof-of-concept: Whisper + open-source VALL-E-X variant running fully locally.
- **Phase 2** – Model quantization and deployment to Snapdragon / Android via Qualcomm AI Hub; measure latency & quality.
- **Phase 3** – Full React Native app with encrypted local store, offline inference, and clean UX for patients & caregivers.
- **Phase 4** – Clinical feedback loop with speech therapists and people living with ALS / post-stroke speech loss.

3. Feasibility and Viability

A. Analysis of the Feasibility of the Idea

- Core models (Whisper small/medium, quantized zero-shot TTS variants) already run on modern smartphones.
- Qualcomm AI Hub and similar toolchains exist specifically for edge deployment and NPU acceleration.
- Short-sample zero-shot voice cloning is a solved research problem (VALL-E, YourTTS, and subsequent open work).
- Storage footprint of a single voice identity is tiny (typically a few MB).
- On-device encryption and secure enclaves are mature on both Android and iOS.

B. Potential Challenges and Risks

- Model size versus latency on lower-end or older devices.
- Maintaining natural prosody and emotional expressiveness over long utterances.
- Regulatory considerations if the product is later positioned as a medical / assistive healthcare device.
- Initial model download size (mitigated by progressive download or pre-loading on specialized hardware).
- User trust and education: patients and families must understand that voice data never leaves the device.

C. Strategies for Overcoming These Challenges

- Aggressive quantization + exclusive NPU offloading for synthesis to keep latency under real-time thresholds.
- Hierarchical / streaming synthesis (generate short phrases first, then continue) for perceived responsiveness.
- Clear non-medical positioning in the first release (“assistive communication tool”) while collecting clinical evidence and preparing for future regulatory pathways.
- Hybrid mode design: pure offline by default; optional encrypted cloud backup only if the user explicitly opts in.
- Transparent UX messaging and open technical documentation about the privacy architecture to build trust.

4. Impact and Benefits

A. Potential Impact on Target Audience

Primary users: People diagnosed with progressive speech-loss conditions (ALS, Parkinson's, post-stroke dysarthria, head & neck cancer) and their families and caregivers.

Secondary users: Elderly individuals who wish to preserve their voice for possible future cognitive or motor decline, and anyone who values long-term voice identity.

EchoLife restores not only the functional ability to communicate but the emotional bond of hearing a loved one's real voice. It reduces isolation, preserves dignity, and keeps personal identity intact even as physical speech declines.

B. Benefits of the Solution

Dimension	Benefit
Social	Preserves dignity, personal identity, and family emotional connection. Reduces social isolation and communication frustration.
Economic	Lowers long-term care communication costs. Reduces dependence on expensive dedicated AAC hardware. Scalable software solution.
Environmental	Fully local processing eliminates continuous cloud inference energy cost. Minimal data-center footprint.
Clinical / Care	Enables clearer, more natural communication with healthcare providers. Supports continuity of care and patient-centered interaction.

Closing Statement for Judges

EchoLife is not another text-to-speech application. It is a privacy-first, lifelong Digital Voice Identity system built entirely on Edge AI. By capturing a person's authentic voice while they can still speak and then letting them use that voice offline forever, EchoLife restores both communication and identity — the two things people lose when speech disappears. The deep technical commitment to on-device encryption, zero-cloud operation, and continuous adaptation is what makes this solution uniquely positioned to change lives.

EchoLife — Code. Create. Conquer.
Preserving voice. Preserving identity. Preserving dignity.